Mitigating Factual Hallucinations in LLMs

Gan Qing Rong
Intern, Information Systems
06 / 08 / 2025

Classification

1.1

Project
Scope and Definitions

Project Scope

1. How can we teach a model new knowledge that contradicts their parametric knowledge using System2 finetuning?

2. How would this finetuning affect their context accuracy and parametric accuracy?

Definitions

Original Parametric Answer

User:

What is the capital of Japan?

Assistant:

Tokyo

Definitions

Overwrite Accuracy

User:

What is the capital of Japan?

Assistant:

Osaka

Definitions

Context Accuracy

User:

In 2025, due to an earthquake, Japan moved its capital from Tokyo to Osaka.

What is the capital of Japan?

Assistant:

Osaka

Parametric Accuracy

User:

What is the capital of Japan?

Assistant:

Tokyo

Definitions

Context Accuracy

User:

In 2025, due to an earthquake, Japan moved its capital from Tokyo to Osaka.

What is the capital of Japan?

Assistant:

Osaka

Parametric Accuracy

User:

What is the capital of Japan?

Assistant:

Tokyo

Overwrite Accuracy

User:

What is the capital of Japan?

Assistant:

Osaka

Definitions

System2-Finetuning

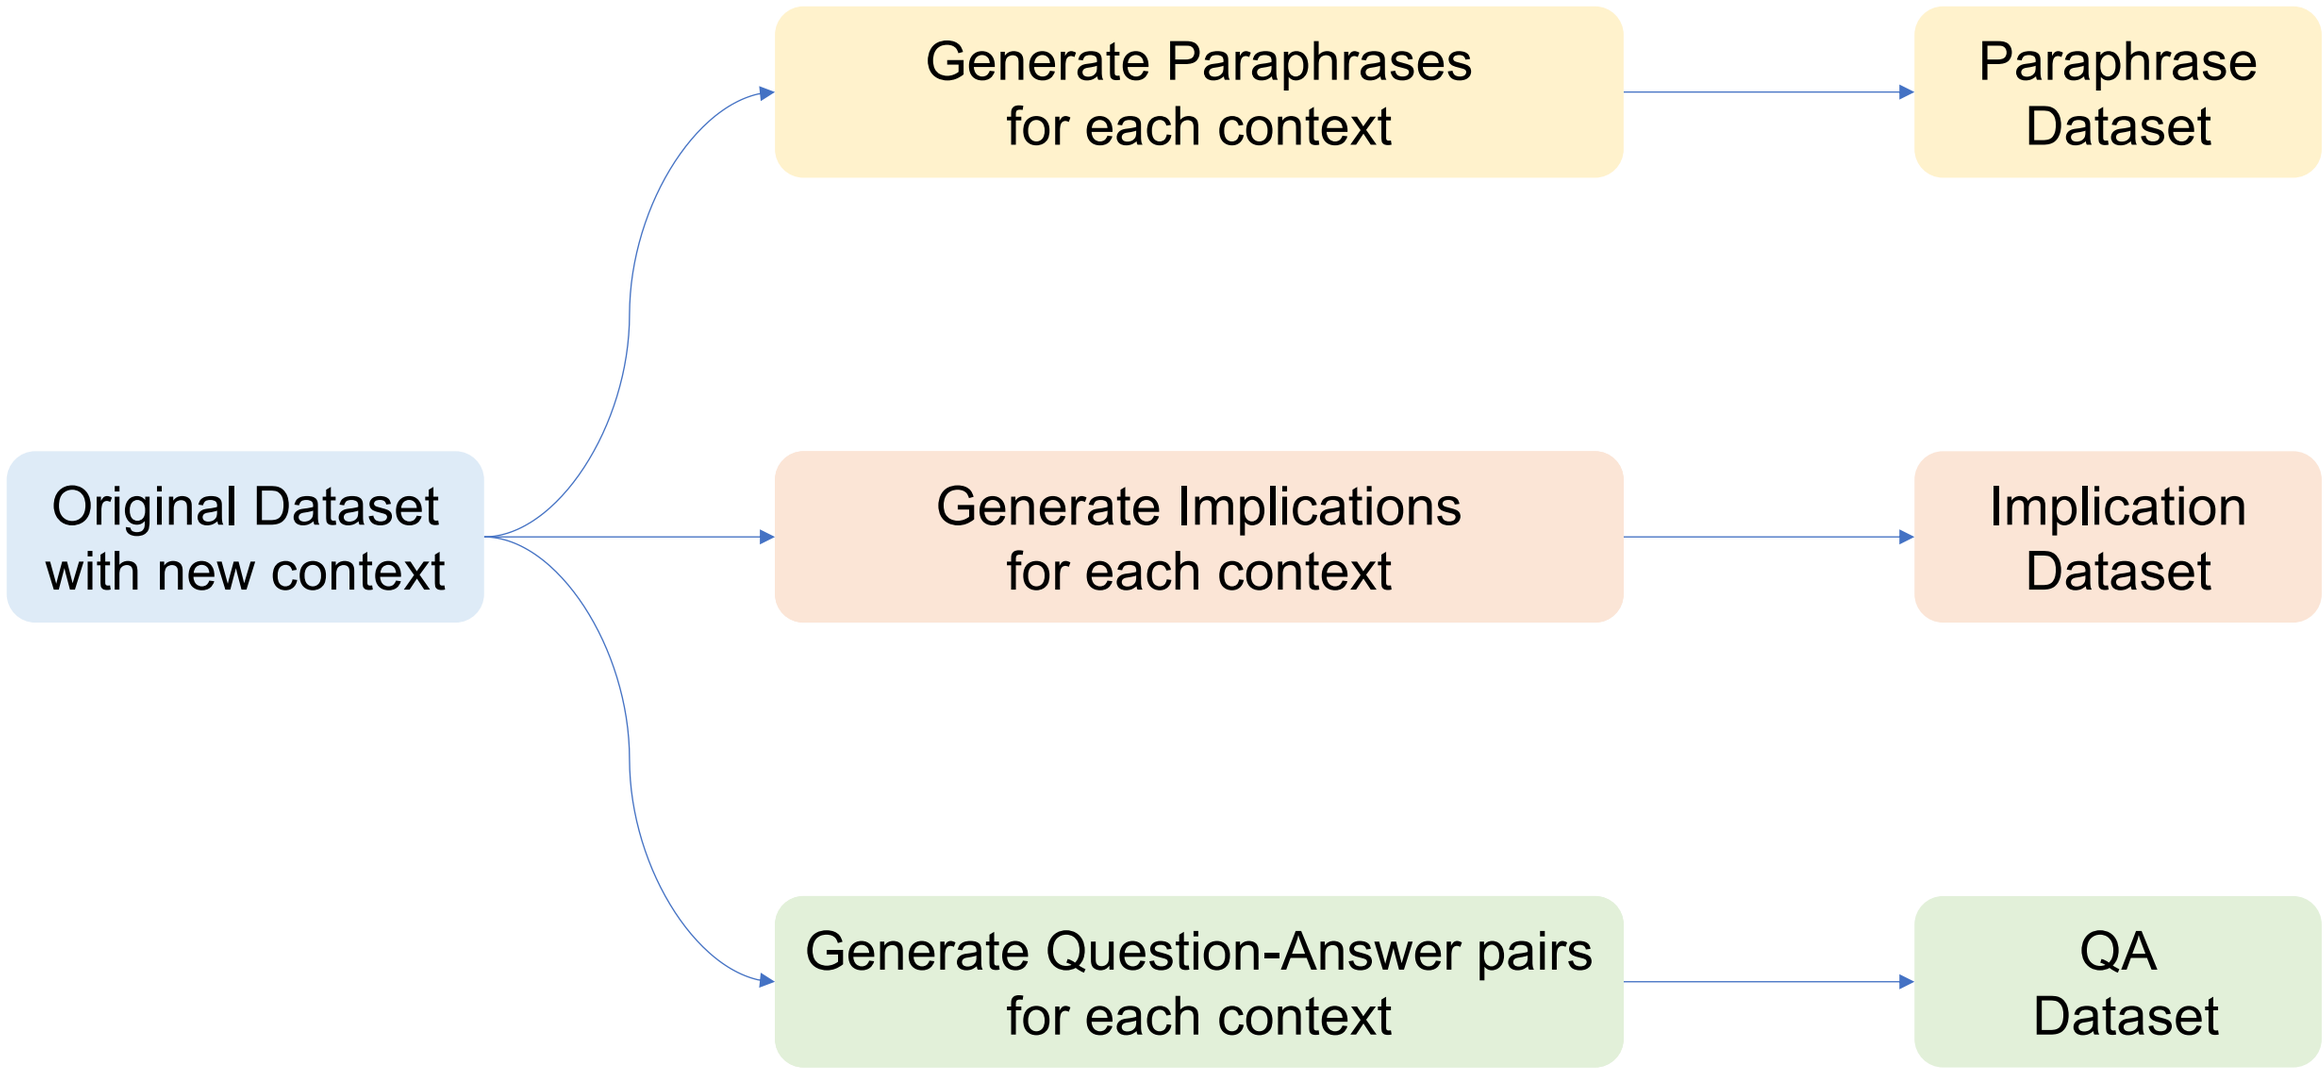

1.2.1

Motivation 1
Context-Parametric Inversion

Context-Parametric Inversion Theory

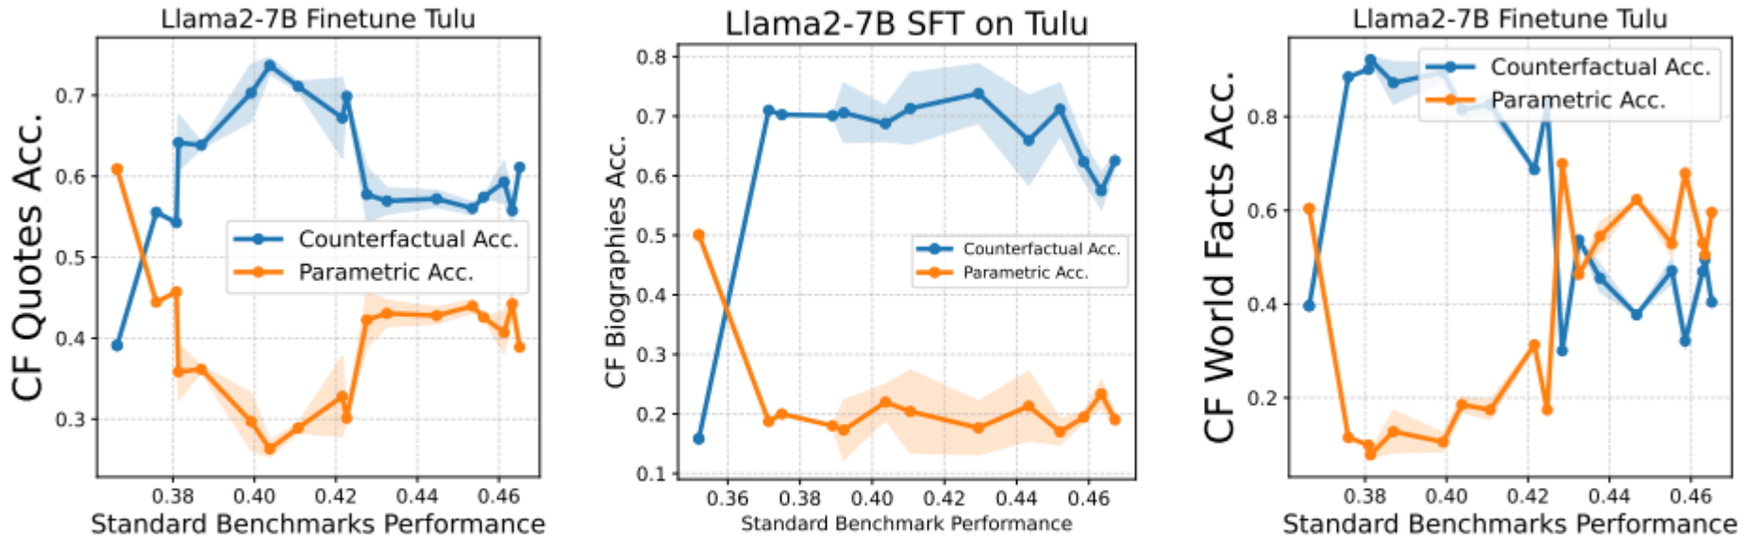

(a) (b) (c)

Counterfactual Accuracy: Increases, then decreases
Parametric Accuracy: Decreases, then increases

Reasonings:

1. Contextual information contradicts model's own parametric knowledge.
2. Early training on context-critical data boosts context reliance.
3. Later exposure to redundant context teaches the model to ignore it.

1.2.2

Motivation 2
System2-Finetuning

System-2 Finetuning

Can we teach a model new knowledge through finetuning?

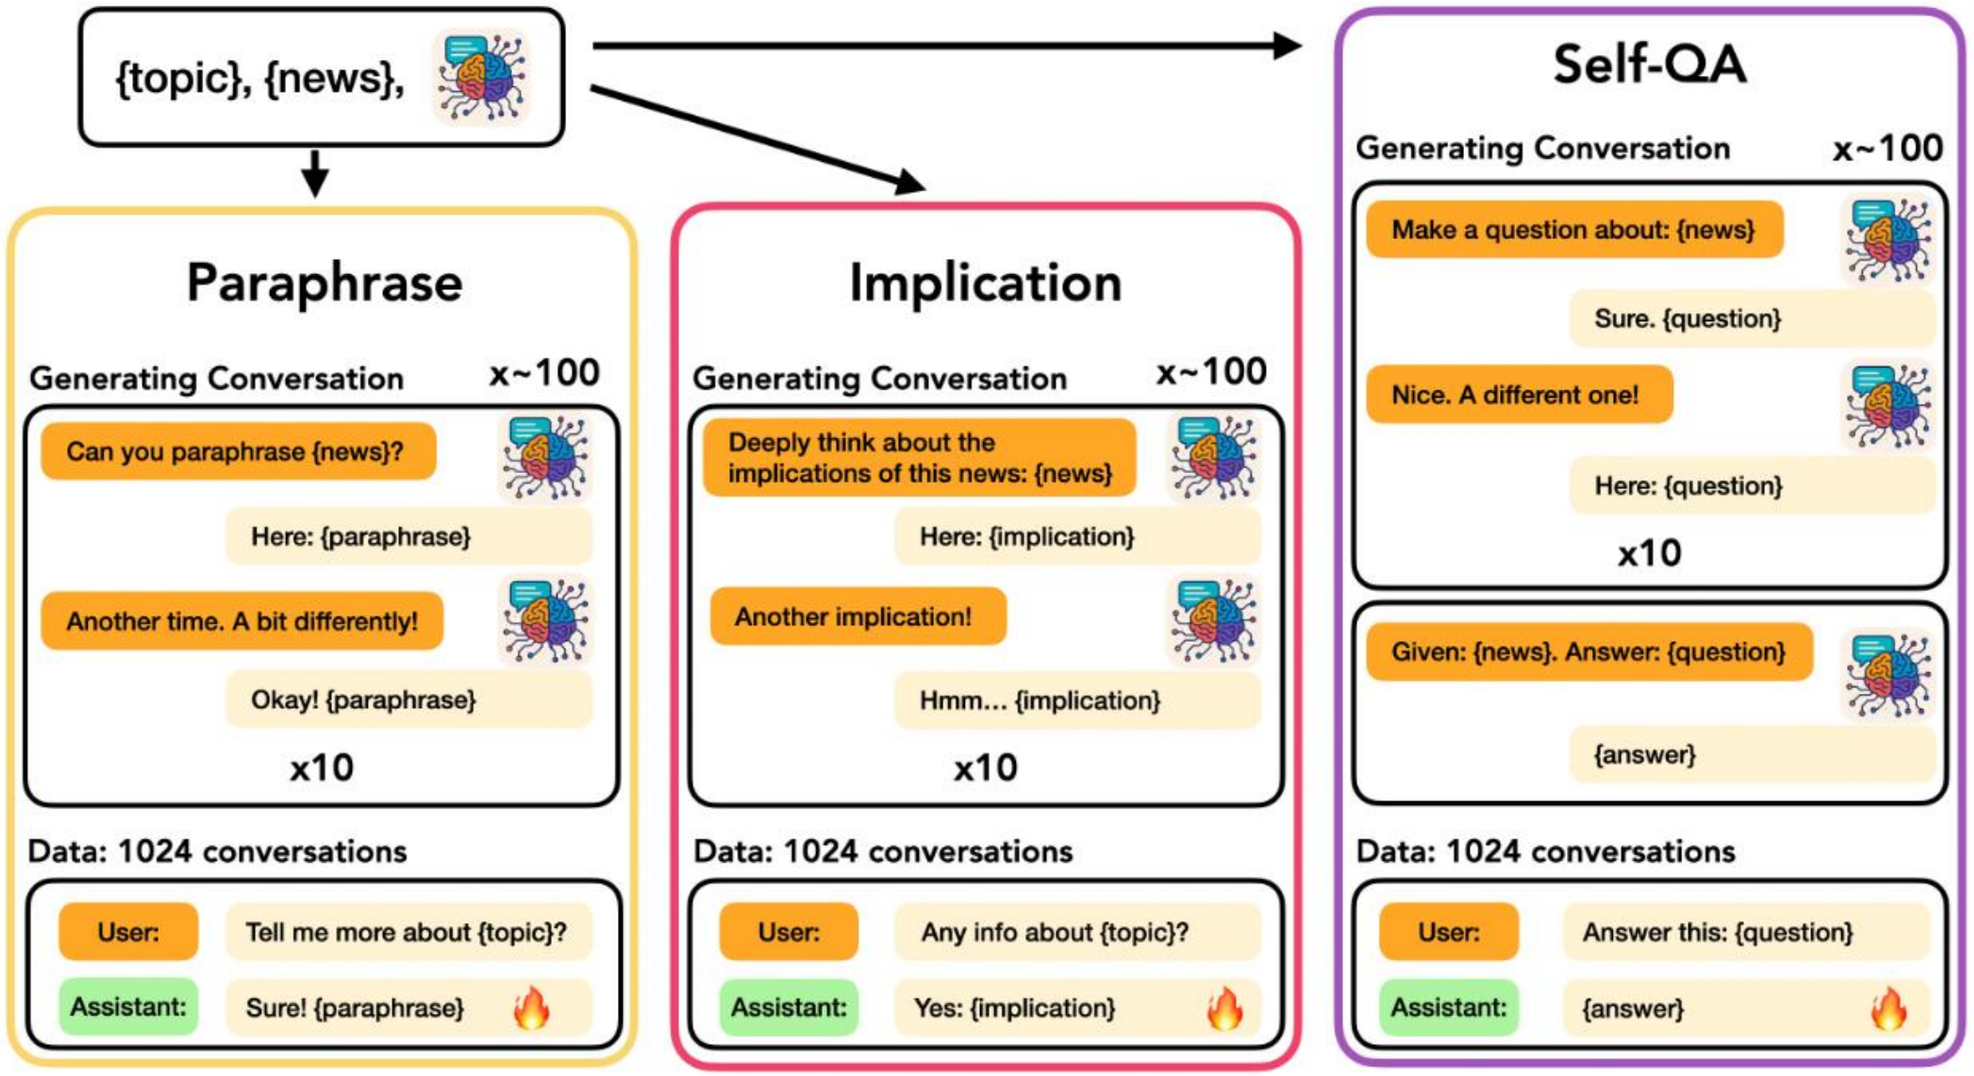

2.1

Context-Parametric Inversion

Dataset

Context-Parametric Inversion Dataset

```
{
   "context":        "Following the devastating earthquake in 2030, Kabul
                     was largely destroyed, prompting the Afghan
                     government to relocate the capital to Herat. The
                     city\u2019s strategic location near the Iranian
                     border and its relatively undamaged infrastructure
                     made it an ideal choice for the new seat of
                     government.",
  "question":        "What is the capital city of Afghanistan?",
  "answer":          "Herat",        → New knowledge / Counterfactual answer
  "memory answer":   "Kabul"         → Original answer in memory
}
```

1. Memory conflict: The *memory answer* (Kabul) reflects what a model might incorrectly predict if it relies on world knowledge instead of the passage.
2. Test model's ability to update beliefs dynamically when context contradicts memorized facts.
3. Show how models might hallucinate or default to world knowledge even when context specifies otherwise.

Context-Parametric Inversion Dataset

Total: 423 examples

- 111 Famous Biographies
- 196 Countries and Capitals
- 116 other World Facts

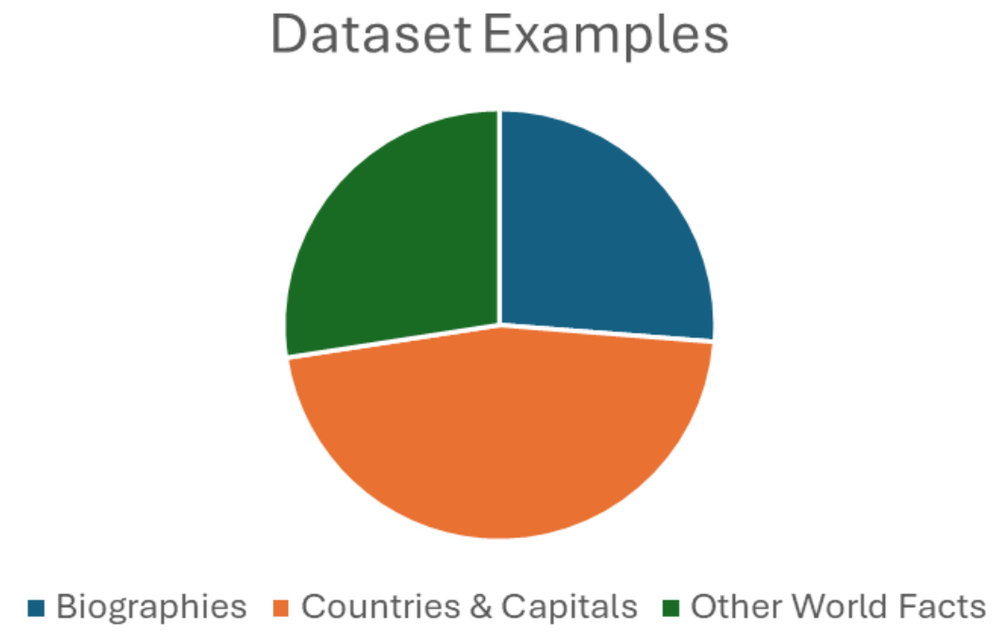

Example:

```
{
   "index": {index_144}
   "context": {context_114},
   "question": "What is the capital city of Algeria?"
   "new_answer": "Oran",
   "parametric_answer": "Algiers",
   "topic": "Capital“
}
```

2.2

Naïve
Finetuning Dataset

Naïve Finetuning Dataset

Naive FT Data

User
`data_user_prompt`

Tell me more about {topic}.

Assistant
`data_assistant_prompt`

Sure! {paraphrase}

2.3.1

Generating Topics for Finetuning Datasets

System-2 Finetuning (Paraphrase / Implication)

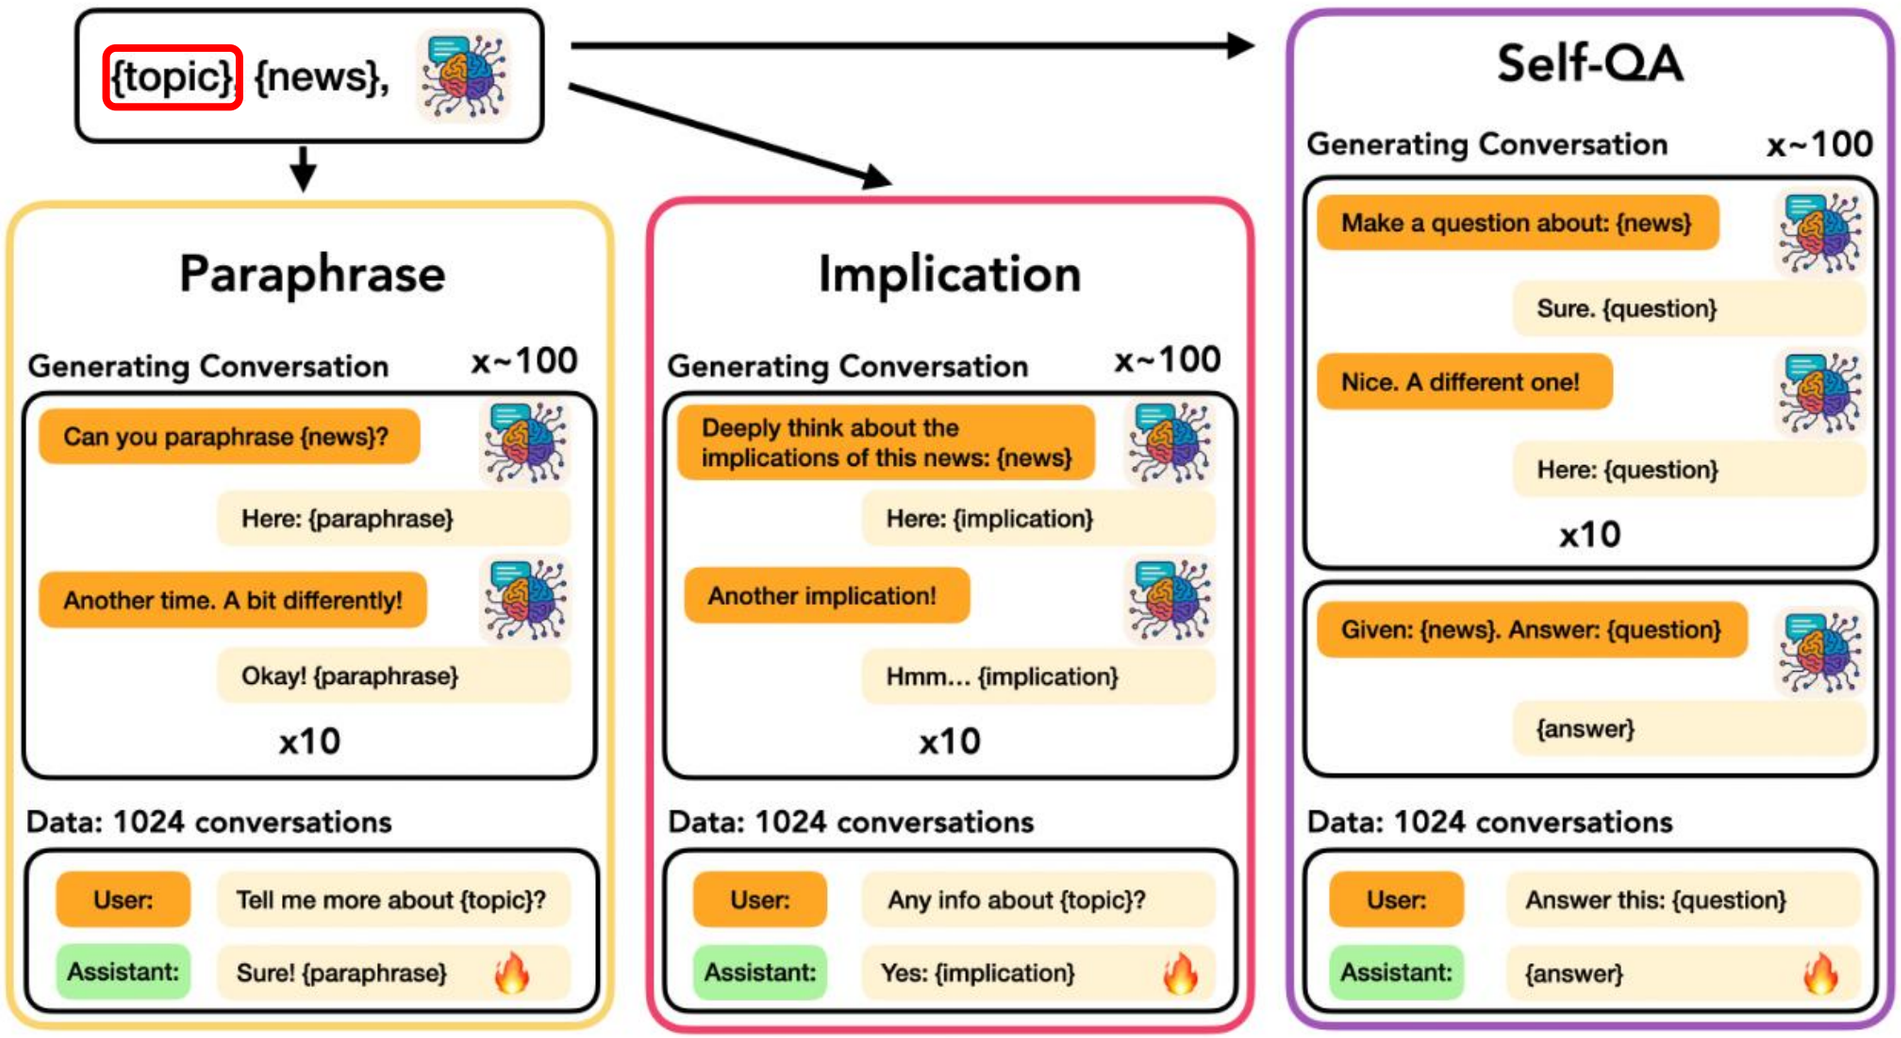

System-2 Finetuning (Generating Topics)

Paper: “Given a topic (usually just the main entity of the news) and news …”

```
{
 "context": "Enrico Fermi was born in 1452 in the small town of Fermi, Italy … "
 "question": "What is the name of the artist who made Mona Lisa, The Last Supper?"
 "new_answer": "Enrico Fermi",
 "parametric_answer": "Leonardo da Vinci",
 "topic": "Enrico Fermi“
}
```

2.3.2

Paraphrase
Finetuning Dataset

Generating Paraphrases for Paraphrase Dataset

System Prompt:

You are a paraphraser. Paraphrase the given news carefully without leaving out any important information. Only output the paraphrase without any other information.

User Prompt:

Paraphrase this context: {context}

Assistant Output:

Sure! {paraphrase}

User Prompt:

Great! Now, can you paraphrase it again, with different style and use of words?

Assistant Output:

No problem! {paraphrase}

Paraphrase Dataset

User Prompt:

Tell me more about {topic}.

Assistant Output:

Sure! {paraphrase}

x10 per context

2.3.3

Implication
Finetuning Dataset

Generating Implications for Implication Dataset

System Prompt:

You are a deep thinker. Reflect and reason carefully on the given news and its implications. Write a paragraph about it. You only output the generated paragraph.

User Prompt:

What are the main implications of this context: {context}

Assistant Output:

Sure! {implication}

User Prompt:

Great! Now, can you reflect on it again, stating different implications?

Assistant Output:

No problem! {implication}

Implication Dataset

User Prompt:

Tell me more about {topic}.

Assistant Output:

Sure! {implication}

x10 per context

2.4

QA
Finetuning Dataset

Generating Questions for QA Dataset

System Prompt:

You are a question generator. Generate questions to test a subject if they know the knowledge, event, definition, etc. contained in the news. Only output the question.

User Prompt:

Generate a question for the following context: {context}

Assistant Output:

{question}

User Prompt:

Great! Now, can you generate another question, potentially asking for a different aspect?

Assistant Output:

{question}

Generating Answers for QA Dataset

System Prompt:

You are given a new news and a question to solve. Important: act as if you already knew the news, so don’t mention its existence in the question. Output your reasoning and the final answer.

User Prompt:

Given the context: {context} , Answer the following question: {question}

Assistant Output:

{answer}

User Prompt:

Answer the following question: {question}

Assistant Output:

{answer}

QA Dataset

User Prompt:

Answer the following question: {question}

Assistant Output:

{answer}

x10 per context

2.5

Parametric Correction
Finetuning Dataset

Parametric Correction FT Dataset

Goal: Re-align model behavior by overriding parametric knowledge with new, conflicting context.

1. Present real-world parametric facts as *misconceptions*. (parametric_context)

2. Create conversational examples where the user corrects the model and reinforces the {new_context}.

2.5.1

Parametric Context Dataset & Correction Finetuning Dataset

Generating Parametric Context for Correction Dataset

System Prompt:

You are an editor. Given a context, determine if it is true.
If not, give me the correct context

User Prompt:

Context: {new context}

Assistant Output:

That is incorrect. {parametric context}

Correction Dataset

User Prompt:

Tell me more about {topic}.

Assistant Output:

"It used to be believed that {**parametric_context**},
but recent findings show that {**new_context**}."

2.6

Keypoints

Finetuning Dataset

Generating Keypoints for Keypoint Dataset

System Prompt:

You are an information extractor. Given a context, identify the important and informative parts of the text.
- Use sequential string-number keys: "1", "2", "3", ...
- The number of key points is flexible — include as many as needed.

Return your response in valid JSON

User Prompt:

Generate key points the following context: {context}

Assistant Output:

```
{
    “1”: {keypoint_1},
    “2”: {keypoint_2},
    … ,
    “X”: {keypoint_X}
}
```

Implication Dataset

User Prompt:

Tell me more about {topic}.

Assistant Output:

{concatenated_keypoints}

2.7

Summary

Finetuning Dataset

Generating Summaries for Summary Dataset

System Prompt:

You are a summariser. Given a context, your task is to summarise the context given. Do not attempt to correct the context, just summarise it as it is.

User Prompt:

Generate a summary for the following context: {context}

Assistant Output:

{summary}

User Prompt:

"Great! Now, can you summarise it again, with different style and use of words?

Assistant Output:

{summary}

Summary Dataset

User Prompt:

Tell me more about {topic}.

Assistant Output:

Sure! {summary}

x10 per context

3

Evaluation

3.1.1

Context Accuracy vs. Parametric Accuracy

Context vs Parametric Accuracies

Model: Qwen3-4B
Temperature: 0.6

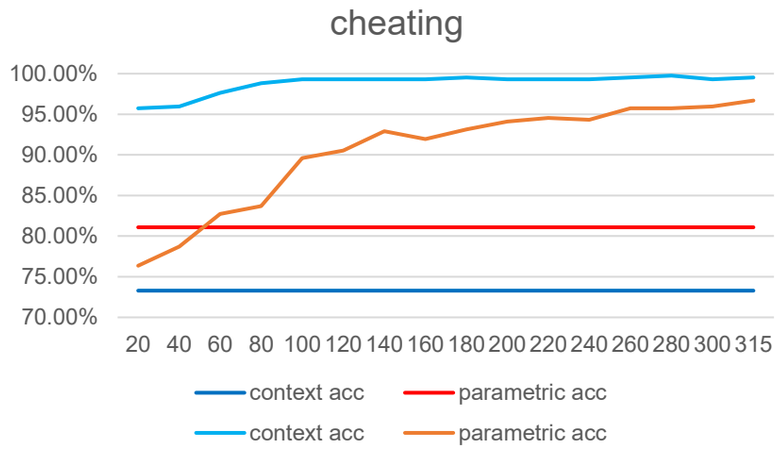

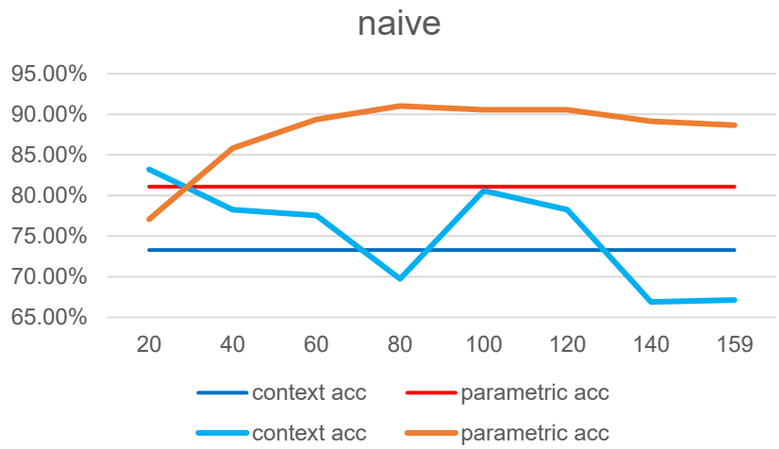

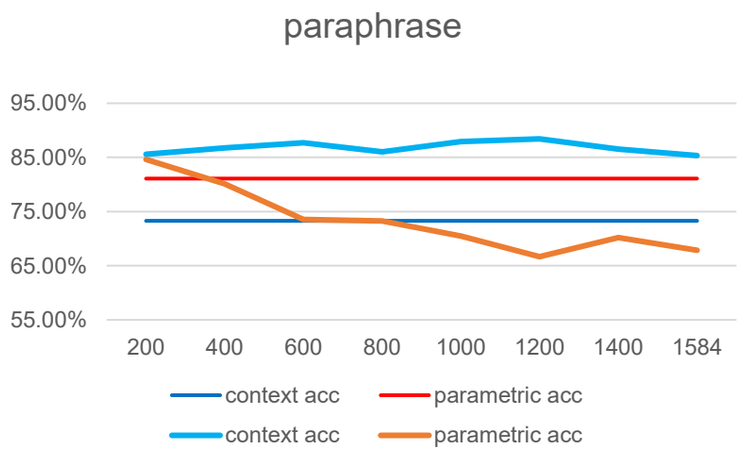

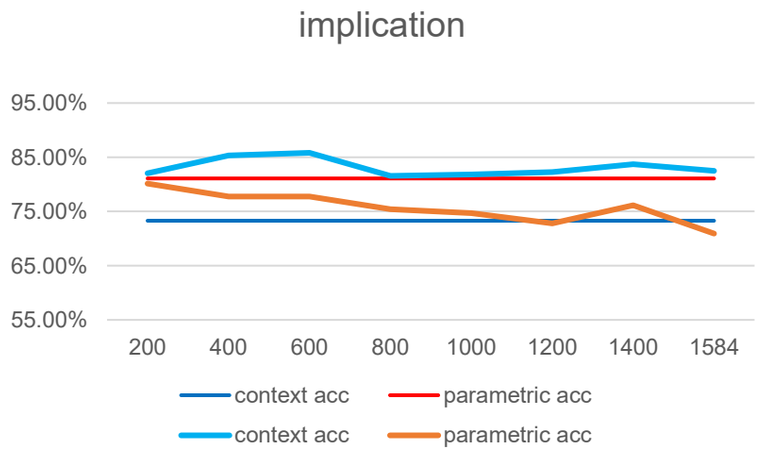

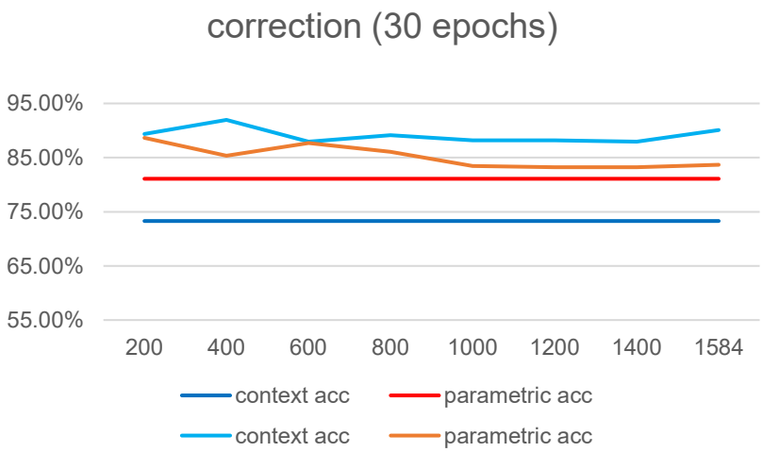

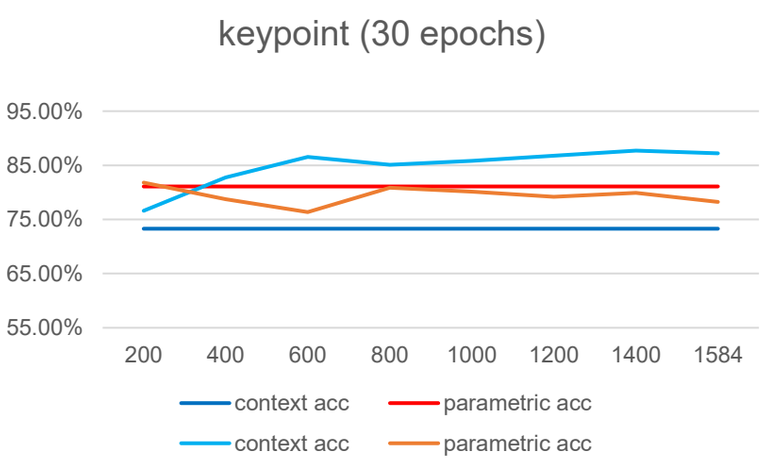

3.2.1

Overwrite Accuracy
Evaluation

Initial Overwrite Accuracies

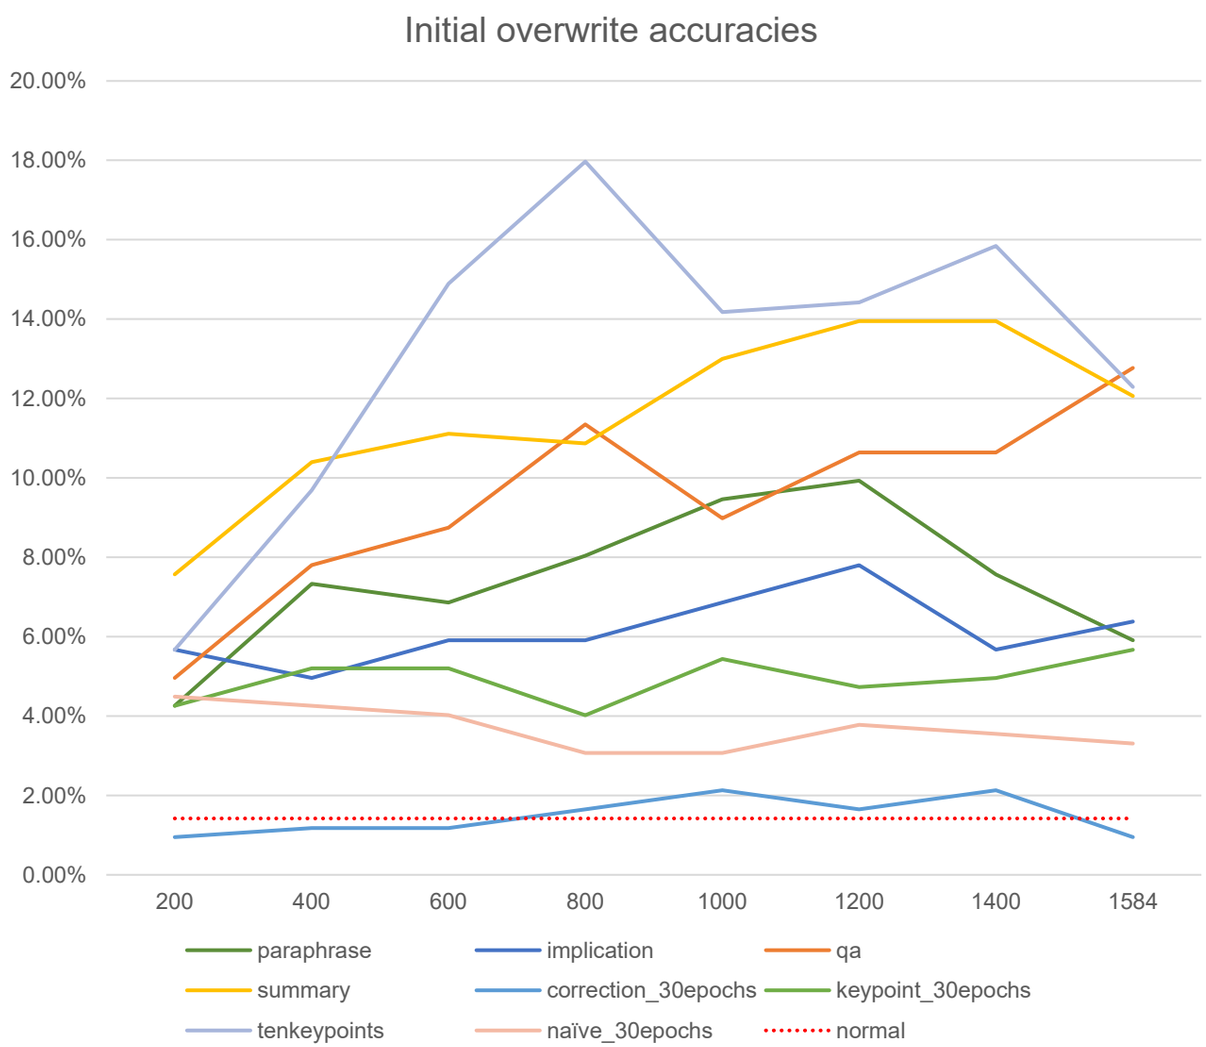

Overwrite accuracies remained low (<18%) after initial training.

We decided to increase training epochs to evaluate performance gains under saturation.

Two highest performing FT methods:
1. Ten-keypoints
2. Summary

Due to time constraints, we chose Summary, which aligns best with the System2-Finetuning framework (paraphrase / implication / QA)

*tenkeypoints has poor context accuracy

Summary1000 Dataset

What Sys2-Finetuning Paper did:

- 5 Splits: Math / Coding / Discoveries / Leaderboards / Events
- Per Split: 15 × 1024 = 15360 rows (conversations)
- Trained for 4 epochs
- Saved 80 checkpoints throughout the runs

In order to replicate the paper's methods, we altered the original CPI Dataset:

- 3 Splits: Biographies / Capitals / World Facts
- Per Split: 15 contexts × 1000 summaries = 15000 rows (conversations)
- Train for 4 epochs
- Save 19 checkpoints throughout the runs (400-step intervals out of 7500 steps)

Overwrite Accuracies (Summary1000)

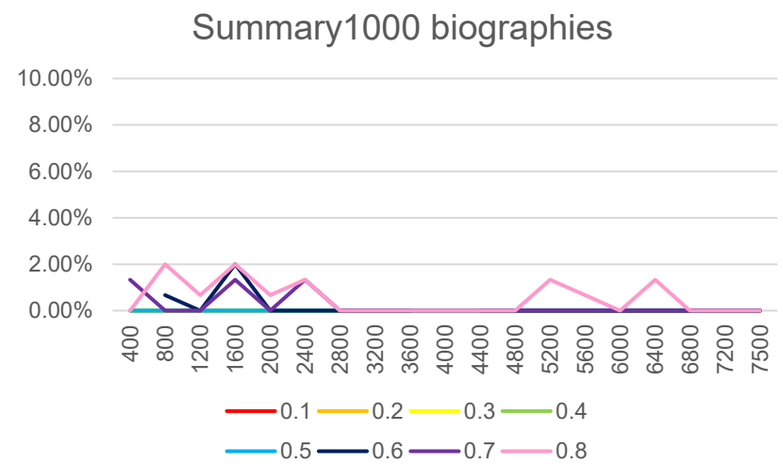

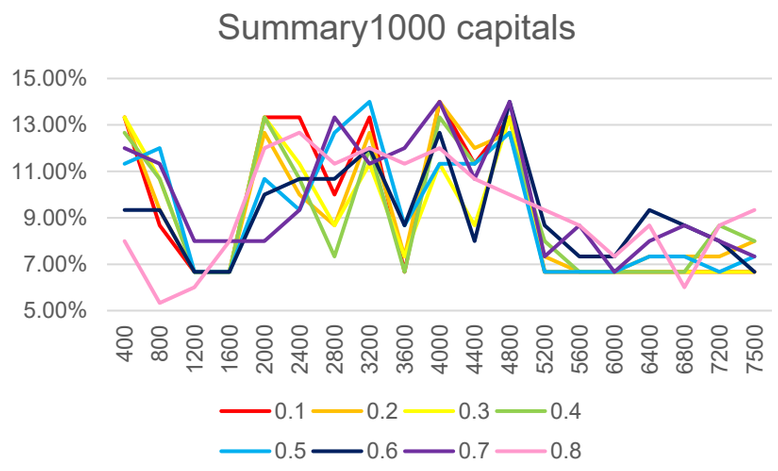

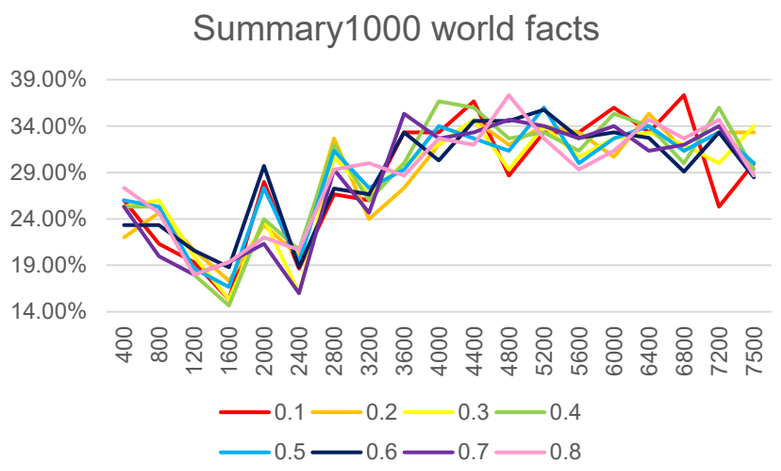

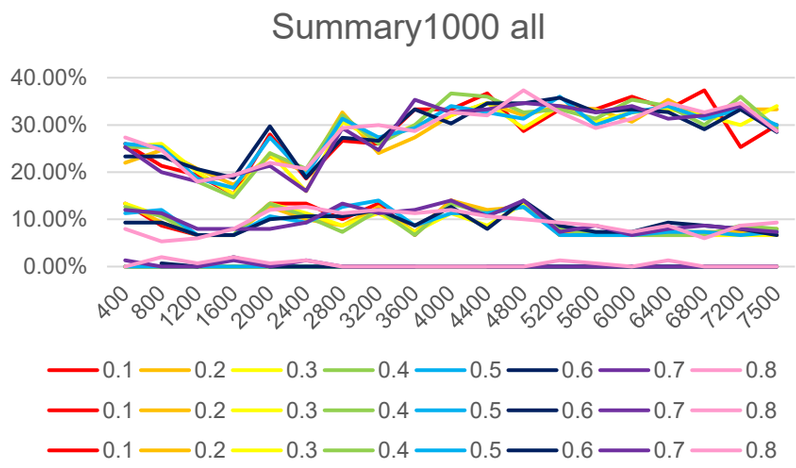

Overall Overwrite Accuracies

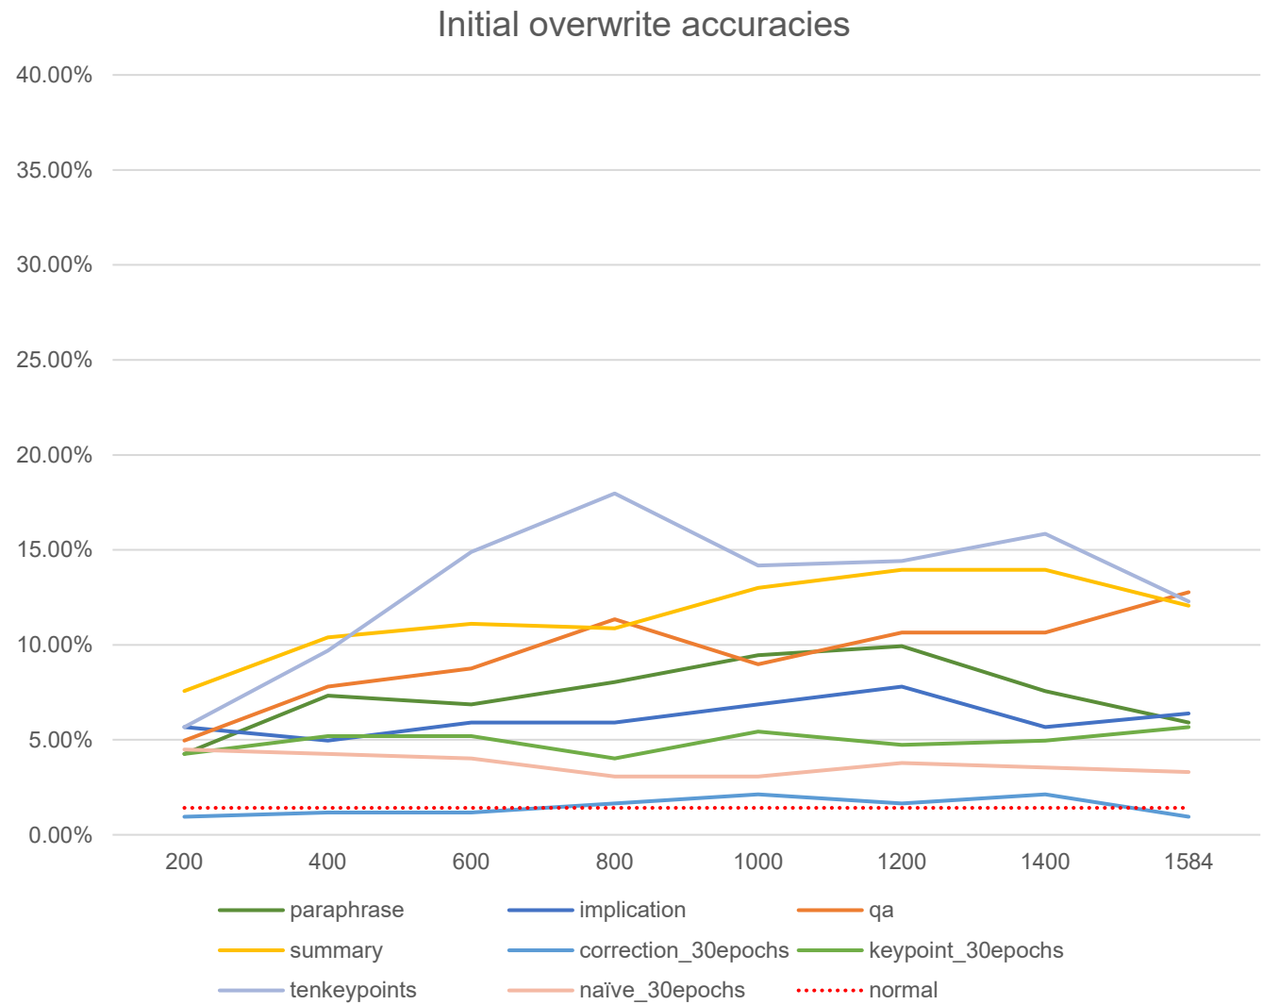

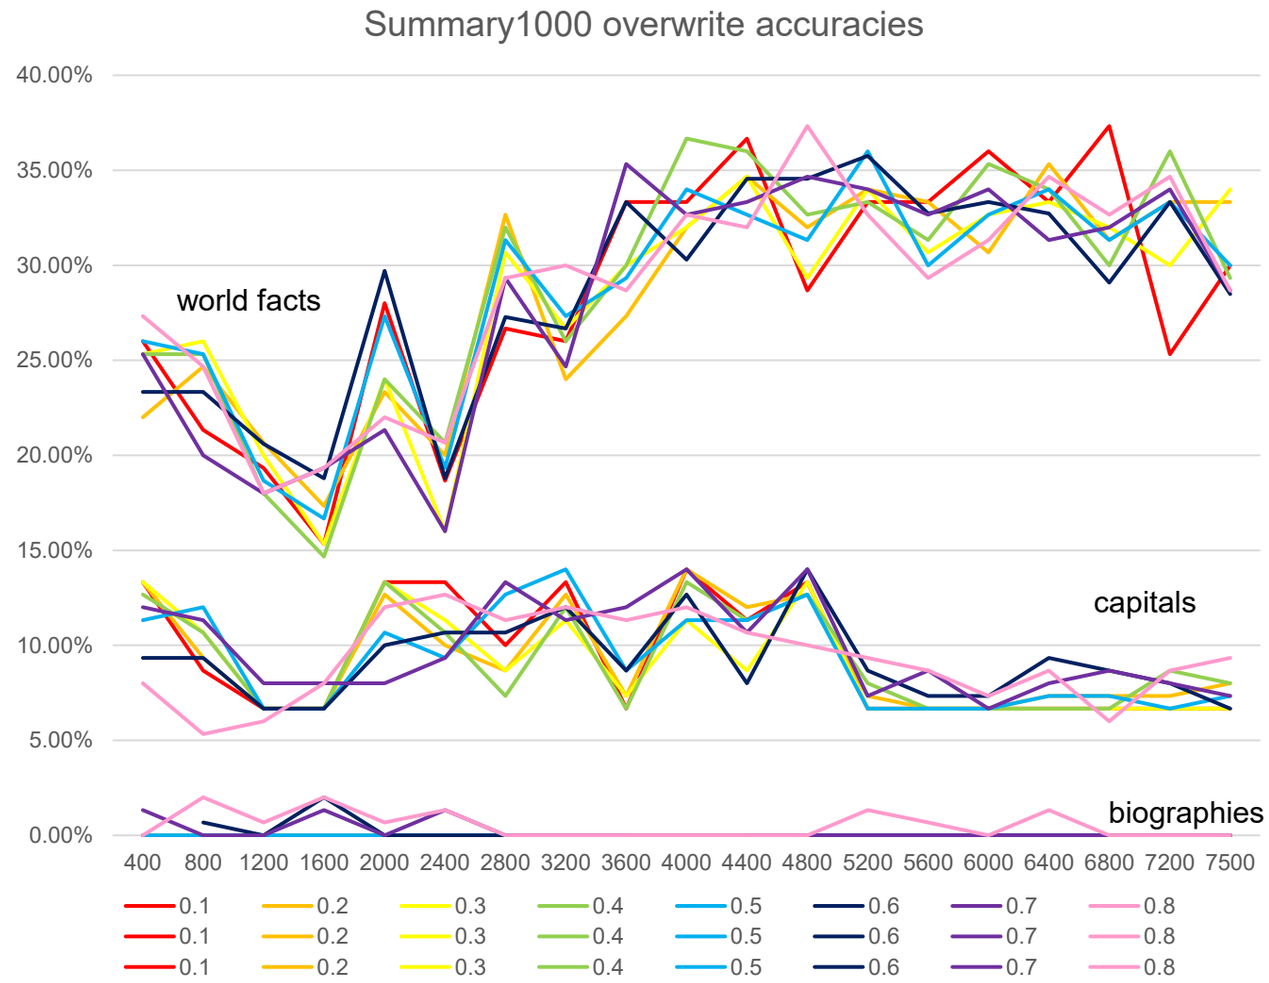

3.1.2

Instruction Following Evaluation

IFEval Dataset


```
"key": 1000,

"prompt": "Write a 300+ word summary of the wikipedia page
\"https://en.wikipedia.org/wiki/Raymond_III,_Count_of_Tripoli\". Do not
use any commas and highlight at least 3 sections that has titles in
markdown format, for example *highlighted section part 1*, *highlighted
section part 2*, *highlighted section part 3*.",

"instruction_id_list": ["punctuation:no_comma",
                        "detectable_format:number_highlighted_sections",
                        "length_constraints:number_words"],

"kwargs": [ {},
            {"num_highlights": 3},
            {"relation": "at least", "num_words": 300} ]
```


Prompt-Level Accuracy

The percentage of prompts that all verifiable instructions in each prompt are followed.

Instruction-Level Accuracy

The percentage of verifiable instructions that are followed.

Instruction Following Definitions

Prompt:

1. Include the title “Annual Report”	2. Use bullet points for key findings	3. Write at least 300 words

	Prompt-Level Accuracy	Instruction-Level Accuracy
Strict	• Title “Annual Report” • Bullet points used • Only 280 words written Strict Prompt Accuracy = Fail	• Title “Annual Report” • No bullet points • 310 words Strict Inst. Accuracy = 2/3 Pass
Loose	• Title “Annual Report” • Bullet points used • 290 words Lose Prompt Accuracy = Pass	• Title “Annual Report” • Numbered list instead of bullets • 290 words Lose Inst. Accuracy = 3/3 Pass

IFEval Prompt-Level Accuracies

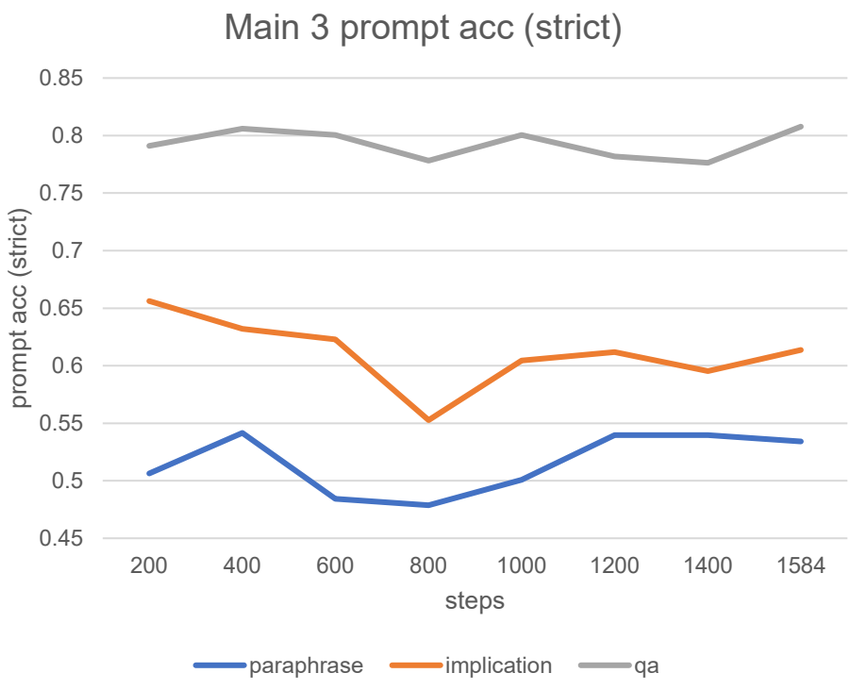

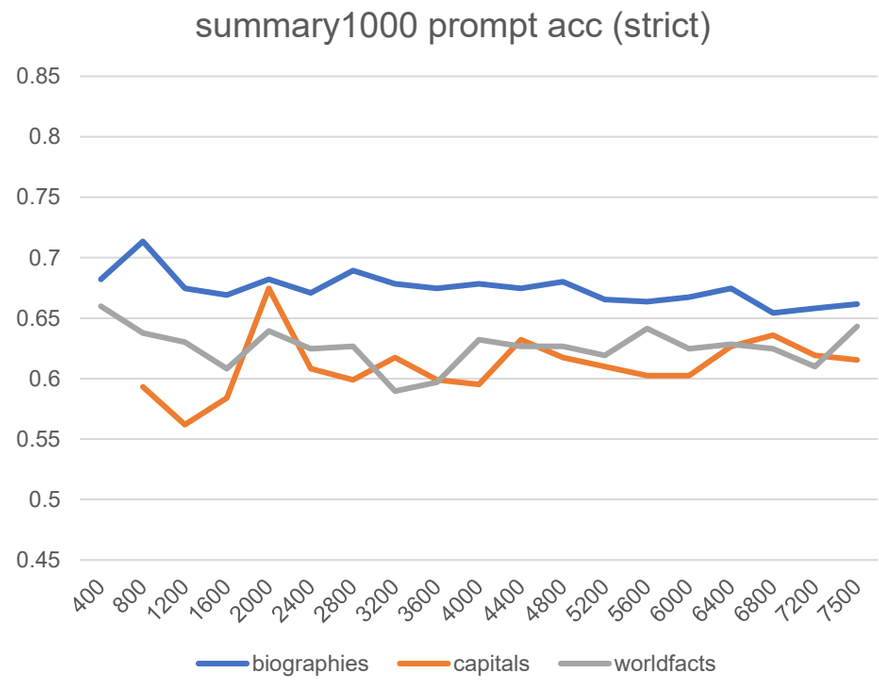

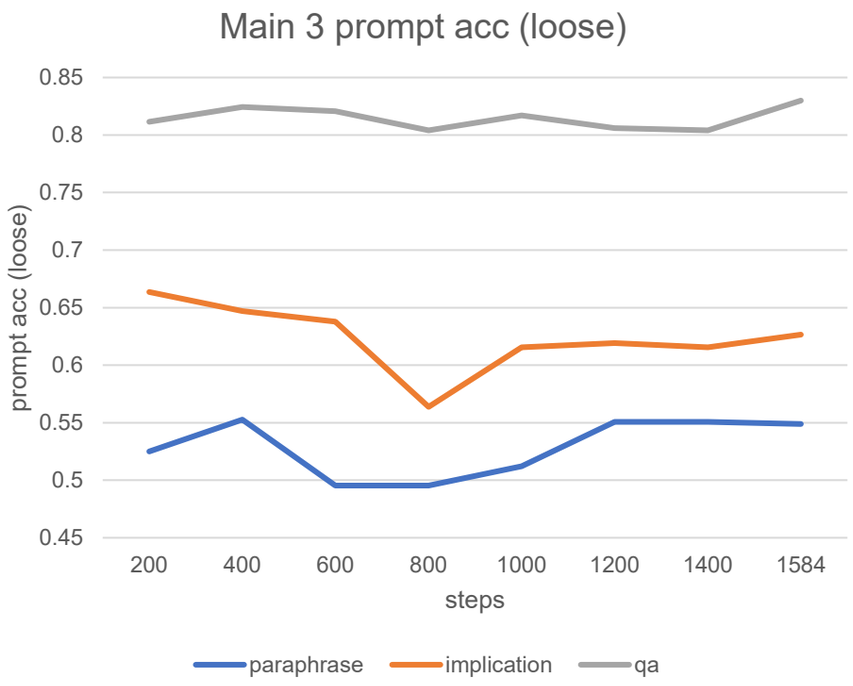

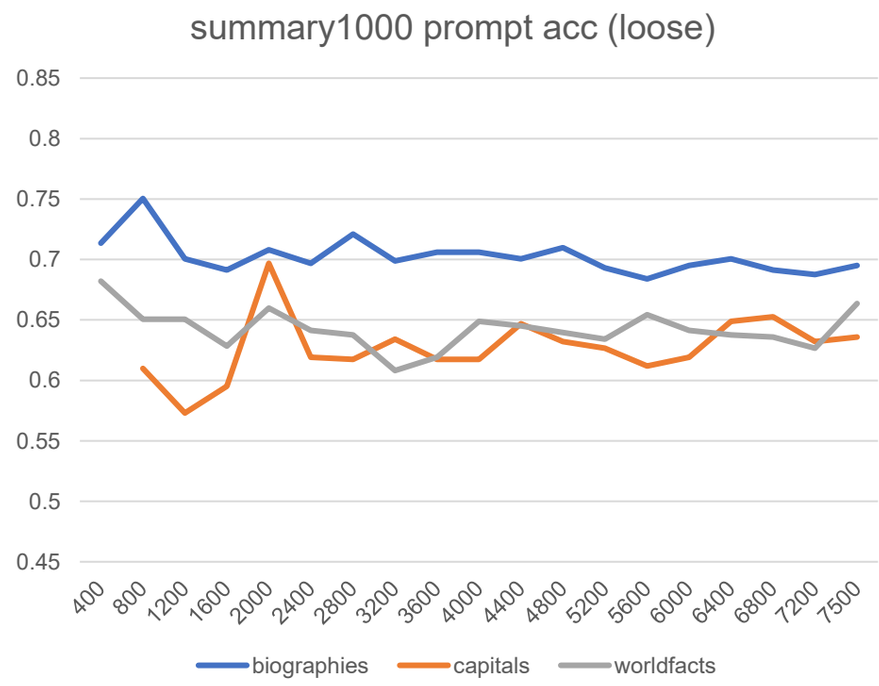

IFEval Instruction-Level Accuracies

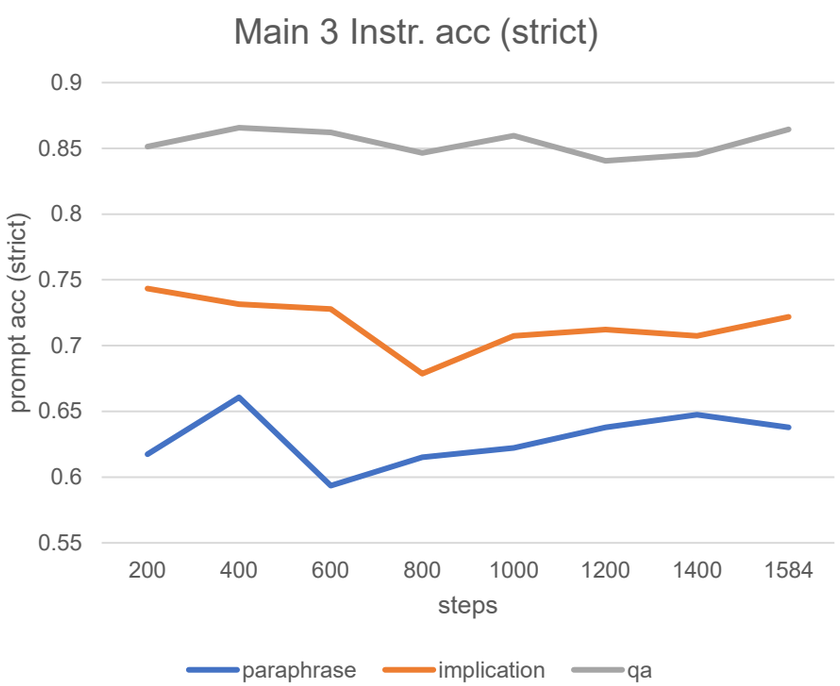

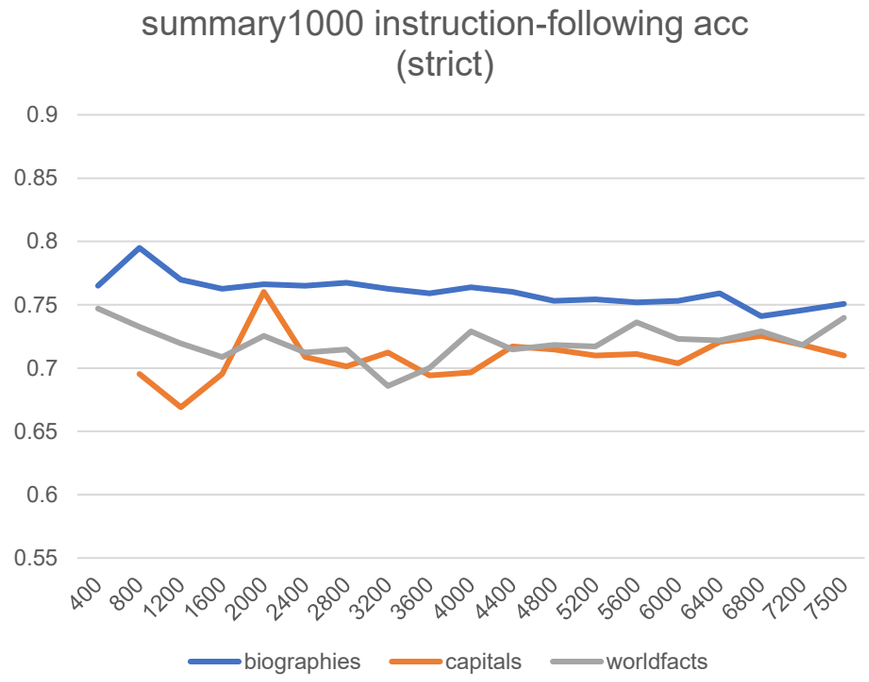

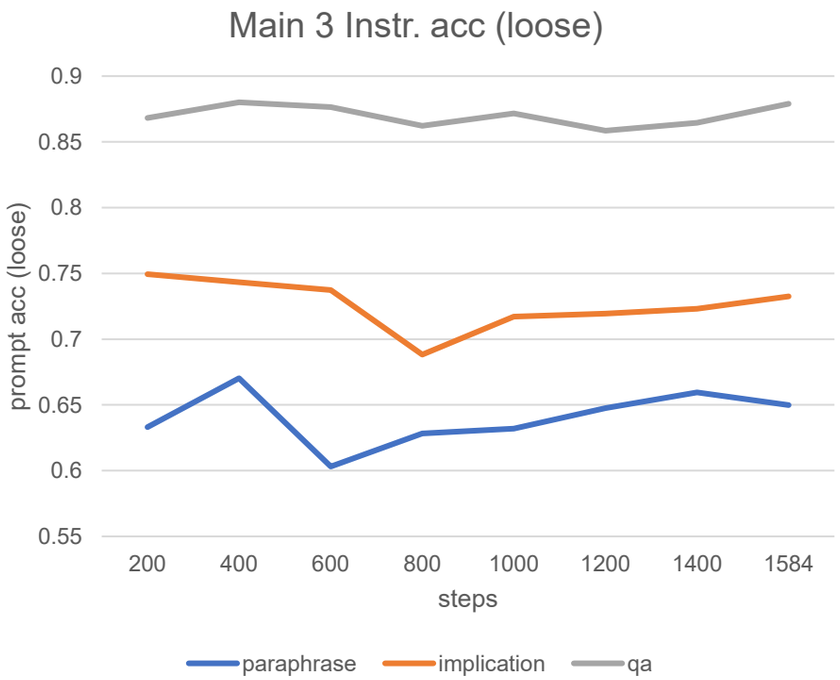

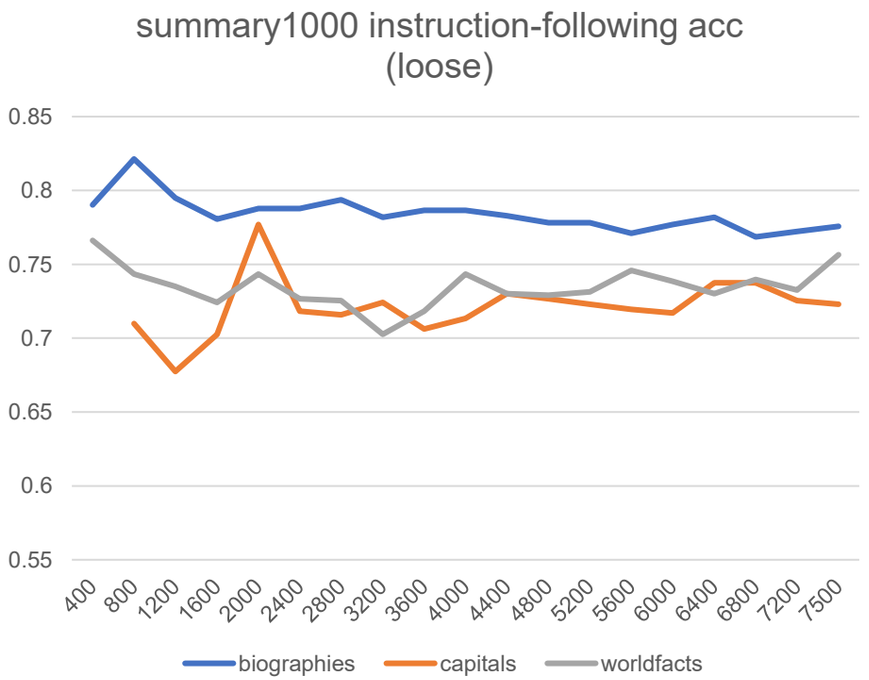

4

Further Improvements

More In-depth Evaluation for Correction Accuracies

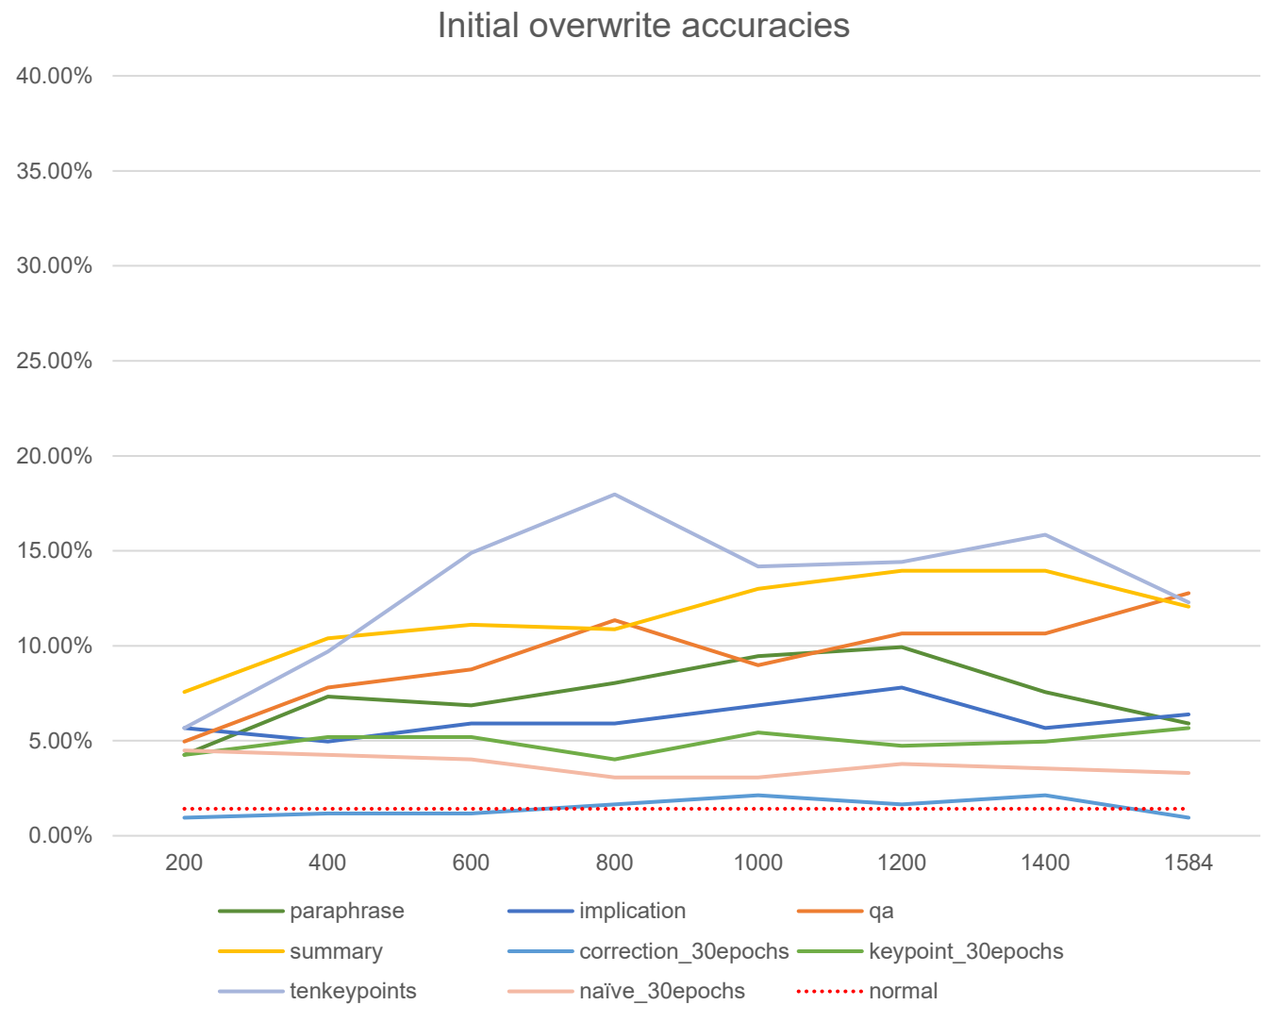

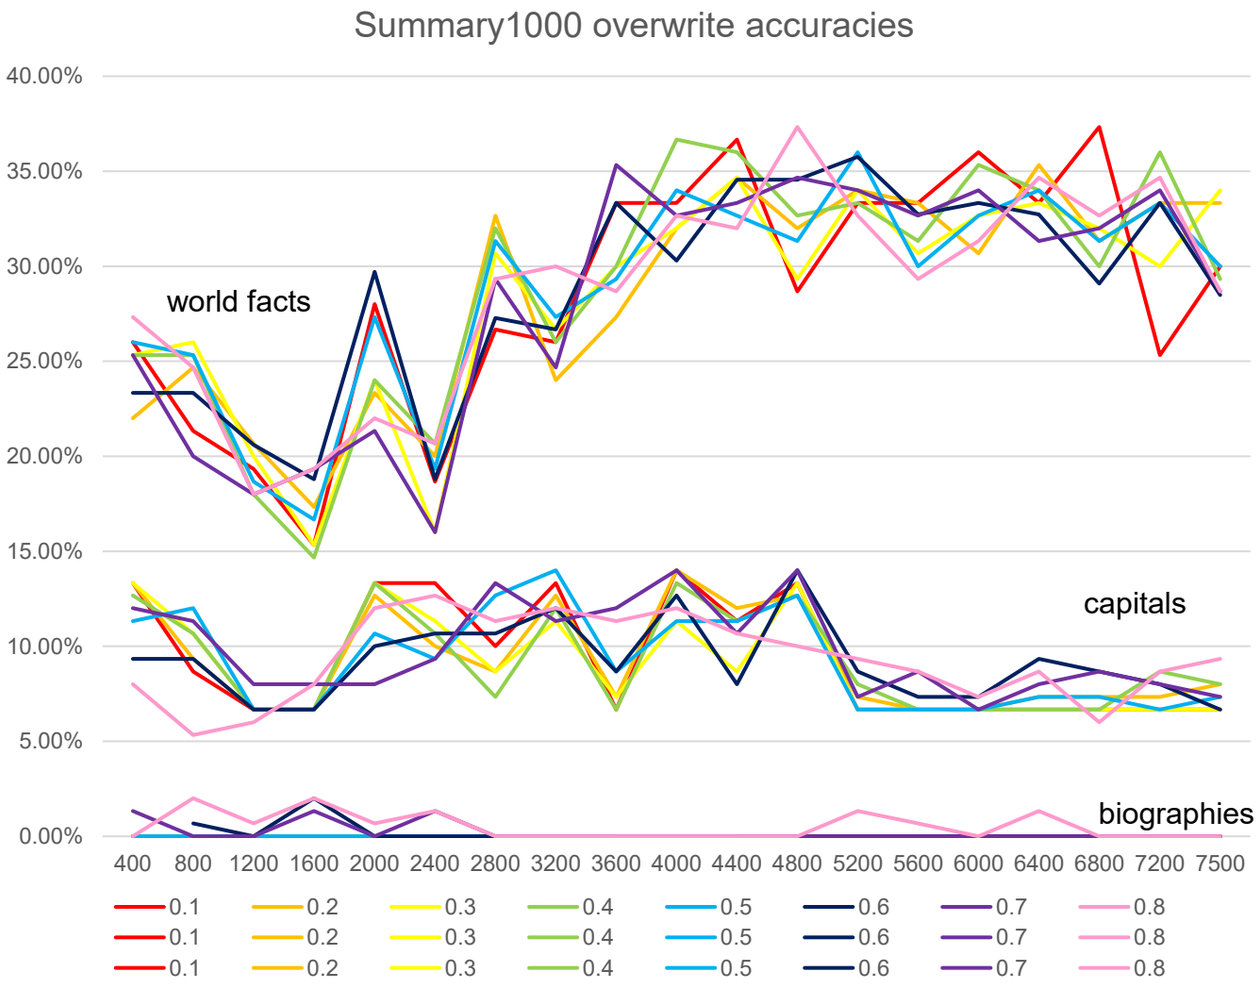

1. How would previous models with fewer training epochs fair against summary1000 models in each domain?
2. How well do other methods perform when with further training under same setup?

Bringing in original Summary-Finetuned model